

Given a bank customer, build a neural network-based classifier that can determine whether they will leave or not in the next 6 months.

Dataset Description: The case study is from an open-source dataset from Kaggle. The dataset contains 10,000 sample points with 14 distinct features such as CustomerId, CreditScore, Geography, Gender, Age, Tenure, Balance, etc. Link to the Kaggle project:

<https://www.kaggle.com/barelydedicated/bank-customer-churn-modeling> Perform following steps:

1. Read the dataset.
2. Distinguish the feature and target set and divide the data set into training and test sets.
3. Normalize the train and test data.
4. Initialize and build the model. Identify the points of improvement and implement the same.
5. Print the accuracy score and confusion matrix.

```
In [46]: import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt #Importing the libraries
```

```
In [47]: df = pd.read_csv("Churn_Modelling.csv")
```

Preprocessing.

```
In [48]: df.head()
```

```
Out[48]:
```

	RowNumber	CustomerId	Surname	CreditScore	Geography	Gender	Age	Tenure	EstimatedSalary
0	1	15634602	Hargrave	619	France	Female	42	2	84504
1	2	15647311	Hill	608	Spain	Female	41	1	85256
2	3	15619304	Onio	502	France	Female	42	8	159360
3	4	15701354	Boni	699	France	Female	39	1	59116
4	5	15737888	Mitchell	850	Spain	Female	43	2	125940

```
In [49]: df.shape
```

Out[49]: (10000, 14)

In [50]: `df.describe()`

Out[50]:

	RowNumber	CustomerId	CreditScore	Age	Tenure	Balance
count	10000.00000	1.000000e+04	10000.000000	10000.000000	10000.000000	10000.000000
mean	5000.50000	1.569094e+07	650.528800	38.921800	5.012800	76485.88928
std	2886.89568	7.193619e+04	96.653299	10.487806	2.892174	62397.40520
min	1.00000	1.556570e+07	350.000000	18.000000	0.000000	0.000000
25%	2500.75000	1.562853e+07	584.000000	32.000000	3.000000	0.000000
50%	5000.50000	1.569074e+07	652.000000	37.000000	5.000000	97198.54000
75%	7500.25000	1.575323e+07	718.000000	44.000000	7.000000	127644.24000
max	10000.00000	1.581569e+07	850.000000	92.000000	10.000000	250898.09000



In [51]: `df.isnull()`

Out[51]:

	RowNumber	CustomerId	Surname	CreditScore	Geography	Gender	Age	Tenure
0	False	False	False	False	False	False	False	False
1	False	False	False	False	False	False	False	False
2	False	False	False	False	False	False	False	False
3	False	False	False	False	False	False	False	False
4	False	False	False	False	False	False	False	False
...
9995	False	False	False	False	False	False	False	False
9996	False	False	False	False	False	False	False	False
9997	False	False	False	False	False	False	False	False
9998	False	False	False	False	False	False	False	False
9999	False	False	False	False	False	False	False	False

10000 rows × 14 columns



In [52]: `df.isnull().sum()`

```
Out[52]: RowNumber      0
         CustomerId     0
         Surname        0
         CreditScore    0
         Geography      0
         Gender         0
         Age            0
         Tenure         0
         Balance        0
         NumOfProducts  0
         HasCrCard      0
         IsActiveMember 0
         EstimatedSalary 0
         Exited         0
         dtype: int64
```

```
In [53]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 10000 entries, 0 to 9999
Data columns (total 14 columns):
 #   Column                Non-Null Count  Dtype
---  ---
 0   RowNumber             10000 non-null  int64
 1   CustomerId            10000 non-null  int64
 2   Surname               10000 non-null  object
 3   CreditScore           10000 non-null  int64
 4   Geography             10000 non-null  object
 5   Gender               10000 non-null  object
 6   Age                  10000 non-null  int64
 7   Tenure               10000 non-null  int64
 8   Balance              10000 non-null  float64
 9   NumOfProducts        10000 non-null  int64
10   HasCrCard            10000 non-null  int64
11   IsActiveMember       10000 non-null  int64
12   EstimatedSalary      10000 non-null  float64
13   Exited               10000 non-null  int64
dtypes: float64(2), int64(9), object(3)
memory usage: 1.1+ MB
```

```
In [54]: df.dtypes
```

```
Out[54]: RowNumber      int64
CustomerId    int64
Surname       object
CreditScore   int64
Geography     object
Gender        object
Age           int64
Tenure        int64
Balance       float64
NumOfProducts int64
HasCrCard     int64
IsActiveMember int64
EstimatedSalary float64
Exited        int64
dtype: object
```

```
In [55]: df.columns
```

```
Out[55]: Index(['RowNumber', 'CustomerId', 'Surname', 'CreditScore', 'Geography',
               'Gender', 'Age', 'Tenure', 'Balance', 'NumOfProducts', 'HasCrCard',
               'IsActiveMember', 'EstimatedSalary', 'Exited'],
              dtype='object')
```

```
In [56]: df = df.drop(['RowNumber', 'Surname', 'CustomerId'], axis=1) #Dropping the unneces
```

```
In [57]: df.head()
```

```
Out[57]:
```


	CreditScore	Geography	Gender	Age	Tenure	Balance	NumOfProducts	HasCrCard
0	619	France	Female	42	2	0.00	1	1
1	608	Spain	Female	41	1	83807.86	1	0
2	502	France	Female	42	8	159660.80	3	1
3	699	France	Female	39	1	0.00	2	0
4	850	Spain	Female	43	2	125510.82	1	1

Visualization

```
In [101... def visualization(x, y, xlabel):
plt.figure(figsize=(10,5))
plt.hist([x, y], color=['red', 'green'], label = ['exit', 'not_exit'])
plt.xlabel(xlabel, fontsize=20)
plt.ylabel("No. of customers", fontsize=20)
plt.legend()
```


```
In [102... df_churn_exited = df[df['Exited']==1]['Tenure']
df_churn_not_exited = df[df['Exited']==0]['Tenure']
```

In [103... visualization(df_churn_exited, df_churn_not_exited, "Tenure")



In [105... df_churn_exited2 = df[df['Exited']==1]['Age']
df_churn_not_exited2 = df[df['Exited']==0]['Age']

In [106... visualization(df_churn_exited2, df_churn_not_exited2, "Age")



Converting the Categorical Variables

```
In [59]: X = df[['CreditScore', 'Gender', 'Age', 'Tenure', 'Balance', 'NumOfProducts', 'HasCrCard',
states = pd.get_dummies(df['Geography'], drop_first = True)
gender = pd.get_dummies(df['Gender'], drop_first = True)
```

```
In [61]: df = pd.concat([df,gender,states], axis = 1)
```

Splitting the training and testing Dataset

```
In [62]: df.head()
```

```
Out[62]:
```

	CreditScore	Geography	Gender	Age	Tenure	Balance	NumOfProducts	HasCrCard
0	619	France	Female	42	2	0.00	1	1
1	608	Spain	Female	41	1	83807.86	1	0
2	502	France	Female	42	8	159660.80	3	1
3	699	France	Female	39	1	0.00	2	0
4	850	Spain	Female	43	2	125510.82	1	1

```
In [63]: X = df[['CreditScore', 'Age', 'Tenure', 'Balance', 'NumOfProducts', 'HasCrCard', 'IsActive']
```

```
In [64]: y = df['Exited']
```

```
In [65]: from sklearn.model_selection import train_test_split
X_train,X_test,y_train,y_test = train_test_split(X,y,test_size = 0.30)
```

Normalizing the values with mean as 0 and Standard Deviation as 1

```
In [66]: from sklearn.preprocessing import StandardScaler
sc = StandardScaler()
```

```
In [67]: X_train = sc.fit_transform(X_train)
X_test = sc.transform(X_test)
```

```
In [68]: X_train
```

```
Out[68]: array([[ 4.56838557e-01, -9.45594735e-01,  1.58341939e-03, ...,
                  9.13181783e-01, -5.81969145e-01, -5.73611200e-01],
                [-2.07591864e-02, -2.77416637e-01,  3.47956411e-01, ...,
                  -1.09507222e+00, -5.81969145e-01,  1.74334114e+00],
                [-1.66115021e-01,  1.82257167e+00, -1.38390855e+00, ...,
                  -1.09507222e+00, -5.81969145e-01, -5.73611200e-01],
                ...,
                [-3.63383654e-01, -4.68324665e-01,  1.73344838e+00, ...,
                  9.13181783e-01, -5.81969145e-01, -5.73611200e-01],
                [ 4.67221117e-01, -1.42286480e+00,  1.38707539e+00, ...,
                  9.13181783e-01, -5.81969145e-01,  1.74334114e+00],
                [-8.82511636e-01,  2.95307447e-01, -6.91162564e-01, ...,
                  9.13181783e-01, -5.81969145e-01, -5.73611200e-01]])
```

```
In [69]: X_test
```

```
Out[69]: array([[ 3.63395520e-01,  1.99853433e-01,  1.58341939e-03, ...,
                  9.13181783e-01, -5.81969145e-01, -5.73611200e-01],
                [-4.15243057e-02,  4.86215475e-01,  1.58341939e-03, ...,
                  -1.09507222e+00, -5.81969145e-01,  1.74334114e+00],
                [-1.87923736e+00, -3.72870651e-01, -1.38390855e+00, ...,
                  9.13181783e-01, -5.81969145e-01, -5.73611200e-01],
                ...,
                [-6.02182526e-01, -5.63778679e-01, -1.73028154e+00, ...,
                  -1.09507222e+00, -5.81969145e-01, -5.73611200e-01],
                [ 1.51585964e+00, -6.59232693e-01,  1.73344838e+00, ...,
                  9.13181783e-01, -5.81969145e-01, -5.73611200e-01],
                [-5.19122049e-01,  1.04399419e-01,  1.73344838e+00, ...,
                  9.13181783e-01, -5.81969145e-01, -5.73611200e-01]])
```

Building the Classifier Model using Keras

```
In [70]: import keras #Keras is the wrapper on the top of tensorflow
         #Can use Tensorflow as well but won't be able to understand the errors initially.
```

```
In [71]: from keras.models import Sequential #To create sequential neural network
         from keras.layers import Dense #To create hidden layers
```

```
In [72]: classifier = Sequential()
```

```
In [74]: #To add the layers
         #Dense helps to construct the neurons
         #Input Dimension means we have 11 features
         #Units is to create the hidden layers
         #Uniform helps to distribute the weight uniformly
         classifier.add(Dense(activation = "relu",input_dim = 11,units = 6,kernel_initialize
```

```
In [75]: classifier.add(Dense(activation = "relu",units = 6,kernel_initializer = "uniform"))
```

```
In [76]: classifier.add(Dense(activation = "sigmoid",units = 1,kernel_initializer = "uniform
```

```
In [77]: classifier.compile(optimizer="adam",loss = 'binary_crossentropy',metrics = ['accuracy'])
```

```
In [79]: classifier.summary() #3 Layers created. 6 neurons in 1st,6neurons in 2nd Layer and
```

Model: "sequential_1"

Layer (type)	Output Shape	Param #
=====		
dense_3 (Dense)	(None, 6)	72

dense_4 (Dense)	(None, 6)	42

dense_5 (Dense)	(None, 1)	7
=====		

Total params: 121

Trainable params: 121

Non-trainable params: 0

```
In [89]: classifier.fit(X_train,y_train,batch_size=10,epochs=50) #Fitting the ANN to trainin
```



```
Epoch 1/50
700/700 [=====] - 0s 674us/step - loss: 0.4293 - accuracy:
0.7947
Epoch 2/50
700/700 [=====] - 0s 647us/step - loss: 0.4239 - accuracy:
0.7947
Epoch 3/50
700/700 [=====] - 0s 657us/step - loss: 0.4203 - accuracy:
0.8067
Epoch 4/50
700/700 [=====] - 0s 664us/step - loss: 0.4167 - accuracy:
0.8260
Epoch 5/50
700/700 [=====] - 0s 674us/step - loss: 0.4153 - accuracy:
0.8287
Epoch 6/50
700/700 [=====] - 0s 653us/step - loss: 0.4137 - accuracy:
0.8310
Epoch 7/50
700/700 [=====] - 0s 658us/step - loss: 0.4125 - accuracy:
0.8317
Epoch 8/50
700/700 [=====] - 1s 842us/step - loss: 0.4116 - accuracy:
0.8306
Epoch 9/50
700/700 [=====] - 0s 671us/step - loss: 0.4103 - accuracy:
0.8331
Epoch 10/50
700/700 [=====] - 0s 682us/step - loss: 0.4100 - accuracy:
0.8326
Epoch 11/50
700/700 [=====] - 0s 690us/step - loss: 0.4093 - accuracy:
0.8337
Epoch 12/50
700/700 [=====] - 0s 688us/step - loss: 0.4087 - accuracy:
0.8339
Epoch 13/50
700/700 [=====] - 0s 675us/step - loss: 0.4081 - accuracy:
0.8341
Epoch 14/50
700/700 [=====] - 1s 722us/step - loss: 0.4071 - accuracy:
0.8331
Epoch 15/50
700/700 [=====] - 1s 811us/step - loss: 0.4065 - accuracy:
0.8341
Epoch 16/50
700/700 [=====] - 0s 711us/step - loss: 0.4056 - accuracy:
0.8356
Epoch 17/50
700/700 [=====] - 0s 702us/step - loss: 0.4046 - accuracy:
0.8366
Epoch 18/50
700/700 [=====] - 0s 688us/step - loss: 0.4035 - accuracy:
0.8343
Epoch 19/50
700/700 [=====] - 1s 715us/step - loss: 0.4024 - accuracy:
```

```
0.8363
Epoch 20/50
700/700 [=====] - 0s 714us/step - loss: 0.4020 - accuracy:
0.8337
Epoch 21/50
700/700 [=====] - 0s 705us/step - loss: 0.4010 - accuracy:
0.8374
Epoch 22/50
700/700 [=====] - 1s 720us/step - loss: 0.4003 - accuracy:
0.8370
Epoch 23/50
700/700 [=====] - 0s 692us/step - loss: 0.3993 - accuracy:
0.8374
Epoch 24/50
700/700 [=====] - 0s 709us/step - loss: 0.3990 - accuracy:
0.8356
Epoch 25/50
700/700 [=====] - 1s 871us/step - loss: 0.3984 - accuracy:
0.8366
Epoch 26/50
700/700 [=====] - 1s 719us/step - loss: 0.3984 - accuracy:
0.8367
Epoch 27/50
700/700 [=====] - 1s 719us/step - loss: 0.3980 - accuracy:
0.8366
Epoch 28/50
700/700 [=====] - 0s 695us/step - loss: 0.3981 - accuracy:
0.8366
Epoch 29/50
700/700 [=====] - 0s 667us/step - loss: 0.3976 - accuracy:
0.8374
Epoch 30/50
700/700 [=====] - 0s 669us/step - loss: 0.3972 - accuracy:
0.8373
Epoch 31/50
700/700 [=====] - 0s 670us/step - loss: 0.3970 - accuracy:
0.8370
Epoch 32/50
700/700 [=====] - 1s 720us/step - loss: 0.3972 - accuracy:
0.8376
Epoch 33/50
700/700 [=====] - 0s 675us/step - loss: 0.3965 - accuracy:
0.8367
Epoch 34/50
700/700 [=====] - 0s 680us/step - loss: 0.3961 - accuracy:
0.8364
Epoch 35/50
700/700 [=====] - 0s 685us/step - loss: 0.3962 - accuracy:
0.8379
Epoch 36/50
700/700 [=====] - 1s 771us/step - loss: 0.3960 - accuracy:
0.8370
Epoch 37/50
700/700 [=====] - 1s 1ms/step - loss: 0.3963 - accuracy: 0.
8366
Epoch 38/50
```

```

700/700 [=====] - 1s 764us/step - loss: 0.3962 - accuracy:
0.8373
Epoch 39/50
700/700 [=====] - 1s 823us/step - loss: 0.3950 - accuracy:
0.8384
Epoch 40/50
700/700 [=====] - 1s 759us/step - loss: 0.3956 - accuracy:
0.8361
Epoch 41/50
700/700 [=====] - 1s 773us/step - loss: 0.3949 - accuracy:
0.8366
Epoch 42/50
700/700 [=====] - 0s 695us/step - loss: 0.3953 - accuracy:
0.8369
Epoch 43/50
700/700 [=====] - 0s 701us/step - loss: 0.3952 - accuracy:
0.8369
Epoch 44/50
700/700 [=====] - 0s 707us/step - loss: 0.3952 - accuracy:
0.8366
Epoch 45/50
700/700 [=====] - 0s 680us/step - loss: 0.3955 - accuracy:
0.8376
Epoch 46/50
700/700 [=====] - 0s 665us/step - loss: 0.3947 - accuracy:
0.8373
Epoch 47/50
700/700 [=====] - 0s 708us/step - loss: 0.3947 - accuracy:
0.8371
Epoch 48/50
700/700 [=====] - 0s 681us/step - loss: 0.3944 - accuracy:
0.8371
Epoch 49/50
700/700 [=====] - 0s 678us/step - loss: 0.3947 - accuracy:
0.8383
Epoch 50/50
700/700 [=====] - 1s 869us/step - loss: 0.3944 - accuracy:
0.8370

```

Out[89]: <tensorflow.python.keras.callbacks.History at 0x1fb1eb93df0>

```
In [90]: y_pred = classifier.predict(X_test)
        y_pred = (y_pred > 0.5) #Predicting the result
```

```
In [97]: from sklearn.metrics import confusion_matrix, accuracy_score, classification_report
```

```
In [92]: cm = confusion_matrix(y_test, y_pred)
```

```
In [93]: cm
```

```
Out[93]: array([[2328,  72],
               [ 425, 175]], dtype=int64)
```


```
In [94]: accuracy = accuracy_score(y_test, y_pred)
```

In [95]: accuracy

Out[95]: 0.8343333333333334

In [98]: plt.figure(figsize = (10,7))
sns.heatmap(cm,annot = True)
plt.xlabel('Predicted')
plt.ylabel('Truth')

Out[98]: Text(69.0, 0.5, 'Truth')



In [100... print(classification_report(y_test,y_pred))

	precision	recall	f1-score	support
0	0.85	0.97	0.90	2400
1	0.71	0.29	0.41	600
accuracy			0.83	3000
macro avg	0.78	0.63	0.66	3000
weighted avg	0.82	0.83	0.81	3000

In []: